

Abdelrahman Elmeligy

AI Workflow Systems Designer · New York, NY · abdelrahmanelmeligy16@gmail.com · [instagram.com/almeligii](https://www.instagram.com/almeligii) (141K) · github.com/abdelrahmanelmeligy16-glitch · [linkedin.com/in/abdelrahman-elmeligy](https://www.linkedin.com/in/abdelrahman-elmeligy)

SUMMARY

I design AI agent systems — the workflow architecture, the human-approval gates, the evaluation loops — and direct AI to build them. In one 90-day stretch (May–Aug 2026) I shipped **five working platforms** on Next.js/TypeScript/Supabase with Claude, MCP, and n8n: a company operating system with 1,158 commits and a real PR workflow, an autonomous revenue operator that runs unattended on scheduled infrastructure and writes its own audit log, and an agent gateway with 289 passing tests. My design signature is **governed autonomy**: systems that act on their own inside boundaries a human sets, log everything they do, and stop where judgment belongs to a person. Every claim on this page is evidence-graded (audit 2026-08-22; methodology and reproduction commands available).

How I build, honestly: AI generates most of the implementation code. I own problem selection, product direction, workflow architecture, constraints, agent orchestration, evaluation criteria, and iteration — the judgment layer. That combination is the skill: one operator directing AI to production-shaped output at team-level speed.

WHAT I'M CAPABLE OF WITH AI

End-to-end system design: full agent workflows from business problem → orchestration → approval gates → execution → self-written audit trail (implemented in 3 production-shaped systems).

Safety-first automation: approval encoded structurally (a "publish" tool that only exists as *publish_approved*); mock-by-default providers so nothing touches live accounts by accident.

Shipping speed: 5 platforms, ~2,000 commits, 90 days, solo — via structured prompt corpora, agent-facing documentation, and multi-machine scheduled infrastructure.

Agent orchestration: primary-brain + advisory-specialist model patterns; a personal memory service (6.8GB index over 1,927 sessions) consulted by other systems before decisions.

Verification discipline: tool certification with named failure modes; 289-test suites reproduced independently; systems that log their own failures with causes.

Translation: built a 141K-follower audience — I can explain what an AI system does to engineers, executives, and customers in their language.

FLAGSHIP PROJECT EXPERIENCE — 2026

KAIRO Intelligence — AI-native company operating system

- Architected a 4,256-file TypeScript platform: 1,158 commits with PR workflow, CI, 62 Supabase migrations, 20 locales incl. RTL (Arabic, Urdu).
- Designed governed operating loops — agents propose, humans approve — with RBAC enforced by 76 behavioral test files (approval gates, GDPR checklist, buyer-readiness).
- Built a tool-certification pipeline grading 529 AI capabilities with named failure modes; prepared an acquisition-grade diligence room (SHA256-verified copy, honest pre-revenue valuation). Status: pre-revenue; deployment currently offline.

Instagram Social Gateway — approval-gated agent access to a live social business

- Architected a TypeScript pnpm monorepo — 18-tool MCP server (bearer auth, rate limiting), Next.js dashboard, DB-polling worker (FOR UPDATE SKIP LOCKED job claiming) + 8 shared packages.
- **289 passing tests** (262 unit + 27 contract), clean typecheck/lint, green production build — independently reproduced in a clean copy.
- Governance in the type system: agents can draft and request approval; publishing exists only for approved items; mock-provider default until live credentials are explicit.

Creator Command (BOLT) — autonomous social-revenue operator, running live

- Operates a 141K-follower account's research → content → distribution → measurement loop on 10 scheduled jobs; the system writes its own 348KB audit log daily, including failures with causes.
- Multi-model orchestration with voice I/O and n8n integration; consults a personal AI memory service as its "advisory brain" before consequential decisions.

THE ABDUL — personal AI memory system

- Ingested 1,927 work sessions into a live 6.8GB SQLite search index served by a dependency-free 152-line Python API, distributed across two machines; consumed by production systems as an advisory brain.

E.L.L.A — creator-business platform

- Designed and shipped a Next.js 16 + Supabase platform (13 public pages, full admin CRUD) with a committed design system — tokens, UI kits, and build documentation written for AI agents as first-class users.

SKILLS — EVIDENCE-TIERED

Verified strong: AI workflow architecture · agent orchestration · prompt & context engineering · human-in-the-loop system design · scheduled automation (queues, retries, health checks) · technical documentation · AI-assisted shipping speed

Demonstrated: evaluation & QA design · Postgres/Supabase schema design · MCP server design · design systems & UX direction · n8n automation · business-model & valuation thinking · multi-machine infrastructure

Developing: production operations · deployment/DevOps · testing depth

Not claimed: ML engineering · hand-coded senior software engineering · production operations at scale — no evidence yet, so not on this page.

ADDITIONAL

Creator & operator of @almeligii (141K Instagram followers). Maintains a public evidence-graded portfolio — every claim carries a verification status and reproduction command: interactive system maps, evidence registry, and per-project proof/limitation records. Journey: Port Said → Kent → Cairo → Los Angeles → New York.